

TV Script Generation

In this project, you'll generate your own [Simpsons](https://en.wikipedia.org/wiki/The_Simpsons) (https://en.wikipedia.org/wiki/The_Simpsons) TV scripts using RNNs. You'll be using part of the [Simpsons dataset](https://www.kaggle.com/wcukierski/the-simpsons-by-the-data) (<https://www.kaggle.com/wcukierski/the-simpsons-by-the-data>) of scripts from 27 seasons. The Neural Network you'll build will generate a new TV script for a scene at [Moe's Tavern](https://simpsonswiki.com/wiki/Moe's_Tavern) (https://simpsonswiki.com/wiki/Moe's_Tavern).

Get the Data

The data is already provided for you. You'll be using a subset of the original dataset. It consists of only the scenes in Moe's Tavern. This doesn't include other versions of the tavern, like "Moe's Cavern", "Flaming Moe's", "Uncle Moe's Family Feed-Bag", etc..

```
In [1]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        import helper

        data_dir = './data/simpsons/moes_tavern_lines.txt'
        text = helper.load_data(data_dir)
        # Ignore notice, since we don't use it for analysing the data
        text = text[81:]
```

Explore the Data

Play around with `view_sentence_range` to view different parts of the data.

```
In [2]: view_sentence_range = (0, 10)

"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

import numpy as np

print('Dataset Stats')
print('Roughly the number of unique words: {}'.format(len({word: None for word
    in text.split()})))
scenes = text.split('\n\n')
print('Number of scenes: {}'.format(len(scenes)))
sentence_count_scene = [scene.count('\n') for scene in scenes]
print('Average number of sentences in each scene: {}'.format(np.average(sen-
    ce_count_scene)))

sentences = [sentence for scene in scenes for sentence in scene.split('\n')]
print('Number of lines: {}'.format(len(sentences)))
word_count_sentence = [len(sentence.split()) for sentence in sentences]
print('Average number of words in each line: {}'.format(np.average(word_count_
    sentence)))

print()
print('The sentences {} to {}'.format(*view_sentence_range))
print('\n'.join(text.split('\n')[view_sentence_range[0]:view_sentence_range[1]
    ]))
```

Dataset Stats

Roughly the number of unique words: 11492

Number of scenes: 262

Average number of sentences in each scene: 15.248091603053435

Number of lines: 4257

Average number of words in each line: 11.50434578341555

The sentences 0 to 10:

Moe_Szyslak: (INTO PHONE) Moe's Tavern. Where the elite meet to drink.

Bart_Simpson: Eh, yeah, hello, is Mike there? Last name, Rotch.

Moe_Szyslak: (INTO PHONE) Hold on, I'll check. (TO BARFLIES) Mike Rotch. Mike Rotch. Hey, has anybody seen Mike Rotch, lately?

Moe_Szyslak: (INTO PHONE) Listen you little puke. One of these days I'm gonna catch you, and I'm gonna carve my name on your back with an ice pick.

Moe_Szyslak: What's the matter Homer? You're not your normal effervescent self.

Homer_Simpson: I got my problems, Moe. Give me another one.

Moe_Szyslak: Homer, hey, you should not drink to forget your problems.

Barney_Gumble: Yeah, you should only drink to enhance your social skills.

Implement Preprocessing Functions

The first thing to do to any dataset is preprocessing. Implement the following preprocessing functions below:

- Lookup Table
- Tokenize Punctuation

Lookup Table

To create a word embedding, you first need to transform the words to ids. In this function, create two dictionaries:

- Dictionary to go from the words to an id, we'll call `vocab_to_int`
- Dictionary to go from the id to word, we'll call `int_to_vocab`

Return these dictionaries in the following tuple (`vocab_to_int`, `int_to_vocab`)

```
In [3]: import numpy as np
        from collections import Counter
        import problem_unittests as tests

        def create_lookup_tables(text):
            """
            Create lookup tables for vocabulary
            :param text: The text of tv scripts split into words
            :return: A tuple of dicts (vocab_to_int, int_to_vocab)
            """

            # TODO: Implement Function
            word_counts = Counter(text)
            sorted_vocab = sorted(word_counts, key=word_counts.get, reverse=True)
            int_to_vocab = {ii: word for ii, word in enumerate(sorted_vocab)}
            vocab_to_int = {word: ii for ii, word in int_to_vocab.items()}

            return vocab_to_int, int_to_vocab

            """
            DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
            """
        tests.test_create_lookup_tables(create_lookup_tables)
```

Tests Passed

Tokenize Punctuation

We'll be splitting the script into a word array using spaces as delimiters. However, punctuations like periods and exclamation marks make it hard for the neural network to distinguish between the word "bye" and "bye!".

Implement the function `token_lookup` to return a dict that will be used to tokenize symbols like "!" into "`||Exclamation_Mark||`". Create a dictionary for the following symbols where the symbol is the key and value is the token:

- Period (.)
- Comma (,)
- Quotation Mark (")
- Semicolon (;)
- Exclamation mark (!)
- Question mark (?)
- Left Parentheses (()
- Right Parentheses ())
- Dash (--)
- Return (\n)

This dictionary will be used to token the symbols and add the delimiter (space) around it. This separates the symbols as it's own word, making it easier for the neural network to predict on the next word. Make sure you don't use a token that could be confused as a word. Instead of using the token "dash", try using something like "`||dash||`".

```
In [4]: def token_lookup():
        """
        Generate a dict to turn punctuation into a token.
        :return: Tokenize dictionary where the key is the punctuation and the value is the token
        """
        # TODO: Implement Function
        token_dict = {
            '.' : "||Period||",
            ',' : "||Comma||",
            '"' : "||Quotation_Mark||",
            ';' : "||Semicolon||",
            '!' : "||Exclamation_Mark||",
            '?' : "||Question_Mark||",
            '(' : "||Left_Parentheses||",
            ')' : "||Right_Parentheses||",
            '--' : "||Dash||",
            '\n' : "||Return||"
        }
        return token_dict

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_tokenize(token_lookup)
```

Tests Passed

Preprocess all the data and save it

Running the code cell below will preprocess all the data and save it to file.

```
In [5]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        # Preprocess Training, Validation, and Testing Data
        helper.preprocess_and_save_data(data_dir, token_lookup, create_lookup_tables)
```

Check Point

This is your first checkpoint. If you ever decide to come back to this notebook or have to restart the notebook, you can start from here. The preprocessed data has been saved to disk.

```
In [6]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

import helper
import numpy as np
import problem_unittests as tests

int_text, vocab_to_int, int_to_vocab, token_dict = helper.load_preprocess()
```

Build the Neural Network

You'll build the components necessary to build a RNN by implementing the following functions below:

- `get_inputs`
- `get_init_cell`
- `get_embed`
- `build_rnn`
- `build_nn`
- `get_batches`

Check the Version of TensorFlow and Access to GPU

```
In [7]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

from distutils.version import LooseVersion
import warnings
import tensorflow as tf

# Check TensorFlow Version
assert LooseVersion(tf.__version__) >= LooseVersion('1.0'), 'Please use Tensor
Flow version 1.0 or newer'
print('TensorFlow Version: {}'.format(tf.__version__))

# Check for a GPU
if not tf.test.gpu_device_name():
    warnings.warn('No GPU found. Please use a GPU to train your neural network.')
else:
    print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))

TensorFlow Version: 1.0.0
Default GPU Device: /gpu:0
```

Input

Implement the `get_inputs()` function to create TF Placeholders for the Neural Network. It should create the following placeholders:

- Input text placeholder named "input" using the [TF Placeholder](https://www.tensorflow.org/api_docs/python/tf/placeholder) (https://www.tensorflow.org/api_docs/python/tf/placeholder) name parameter.
- Targets placeholder
- Learning Rate placeholder

Return the placeholders in the following tuple (Input, Targets, LearningRate)

```
In [8]: def get_inputs():
        """
        Create TF Placeholders for input, targets, and Learning rate.
        :return: Tuple (input, targets, learning rate)
        """
        # TODO: Implement Function
        input = tf.placeholder(tf.int32, [None, None], name='input')
        targets = tf.placeholder(tf.int32, [None, None], name='targets')
        learning_rate = tf.placeholder(tf.float32, name='learning_rate')
        return input, targets, learning_rate

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_inputs(get_inputs)
```

Tests Passed

Build RNN Cell and Initialize

Stack one or more [BasicLSTMCells](https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/BasicLSTMCell) (https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/BasicLSTMCell) in a [MultiRNNCell](https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell) (https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell).

- The Rnn size should be set using `rnn_size`
- Initialize Cell State using the MultiRNNCell's `zero_state()` (https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell#zero_state) function
 - Apply the name "initial_state" to the initial state using `tf.identity()` (https://www.tensorflow.org/api_docs/python/tf/identity).

Return the cell and initial state in the following tuple (Cell, InitialState)

```
In [9]: def get_init_cell(batch_size, rnn_size):
        """
        Create an RNN Cell and initialize it.
        :param batch_size: Size of batches
        :param rnn_size: Size of RNNs
        :return: Tuple (cell, initialize state)
        """

        # TODO: Implement Function
        lstm = tf.contrib.rnn.BasicLSTMCell(rnn_size)

        # Add dropout to the cell
        drop = tf.contrib.rnn.DropoutWrapper(lstm, 0.5)

        # multiple LSTM layers
        #cell = tf.contrib.rnn.MultiRNNCell([lstm] * 1)

        cell = tf.contrib.rnn.MultiRNNCell([drop] * 1)

        # initialize LSTM cell state and identify as 'initial_state'
        initial_state = tf.identity(cell.zero_state(batch_size, tf.float32), name='initial_state')

        return cell, initial_state

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_init_cell(get_init_cell)
```

Tests Passed

Word Embedding

Apply embedding to input_data using TensorFlow. Return the embedded sequence.


```
In [10]: def get_embed(input_data, vocab_size, embed_dim):
        """
        Create embedding for <input_data>.
        :param input_data: TF placeholder for text input.
        :param vocab_size: Number of words in vocabulary.
        :param embed_dim: Number of embedding dimensions
        :return: Embedded input.
        """

        # TODO: Implement Function
        embedding = tf.Variable(tf.random_uniform((vocab_size, embed_dim), -1,
1))
        embed = tf.nn.embedding_lookup(embedding, input_data)
        return embed

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_embed(get_embed)
```

Tests Passed

Build RNN

You created a RNN Cell in the `get_init_cell()` function. Time to use the cell to create a RNN.

- Build the RNN using the `tf.nn.dynamic_rnn()`.
(https://www.tensorflow.org/api_docs/python/tf/nn/dynamic_rnn)
 - Apply the name "final_state" to the final state using `tf.identity()`.
(https://www.tensorflow.org/api_docs/python/tf/identity)

Return the outputs and final_state state in the following tuple (Outputs, FinalState)

```
In [11]: def build_rnn(cell, inputs):
        """
        Create a RNN using a RNN Cell
        :param cell: RNN Cell
        :param inputs: Input text data
        :return: Tuple (Outputs, Final State)
        """

        # TODO: Implement Function
        outputs, final = tf.nn.dynamic_rnn(cell, inputs, dtype=tf.float32)
        final_state = tf.identity(input=final, name="final_state")
        return outputs, final_state

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_build_rnn(build_rnn)
```

Tests Passed

Build the Neural Network

Apply the functions you implemented above to:

- Apply embedding to input_data using your get_embed(input_data, vocab_size, embed_dim) function.
- Build RNN using cell and your build_rnn(cell, inputs) function.
- Apply a fully connected layer with a linear activation and vocab_size as the number of outputs.

Return the logits and final state in the following tuple (Logits, FinalState)

```
In [12]: def build_nn(cell, rnn_size, input_data, vocab_size, embed_dim):
        """
        Build part of the neural network
        :param cell: RNN cell
        :param rnn_size: Size of rnn
        :param input_data: Input data
        :param vocab_size: Vocabulary size
        :param embed_dim: Number of embedding dimensions
        :return: Tuple (Logits, FinalState)
        """

        # TODO: Implement Function
        emb = get_embed(input_data, vocab_size, embed_dim)
        outputs, final_state = build_rnn(cell, emb)
        logits = tf.contrib.layers.fully_connected(outputs, vocab_size, activation
        _fn=None)
        return logits, final_state

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_build_nn(build_nn)
```

Tests Passed

Batches

Implement `get_batches` to create batches of input and targets using `int_text`. The batches should be a Numpy array with the shape (number of batches, 2, batch size, sequence length). Each batch contains two elements:

- The first element is a single batch of **input** with the shape [batch size, sequence length]
- The second element is a single batch of **targets** with the shape [batch size, sequence length]

If you can't fill the last batch with enough data, drop the last batch.

For example, `get_batches([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20], 3, 2)` would return a Numpy array of the following:

```
[
  # First Batch
  [
    # Batch of Input
    [[ 1  2], [ 7  8], [13 14]]
    # Batch of targets
    [[ 2  3], [ 8  9], [14 15]]
  ]

  # Second Batch
  [
    # Batch of Input
    [[ 3  4], [ 9 10], [15 16]]
    # Batch of targets
    [[ 4  5], [10 11], [16 17]]
  ]

  # Third Batch
  [
    # Batch of Input
    [[ 5  6], [11 12], [17 18]]
    # Batch of targets
    [[ 6  7], [12 13], [18  1]]
  ]
]
```

Notice that the last target value in the last batch is the first input value of the first batch. In this case, 1. This is a common technique used when creating sequence batches, although it is rather unintuitive.

```
In [13]: def get_batches(int_text, batch_size, seq_length):
        """
        Return batches of input and target
        :param int_text: Text with the words replaced by their ids
        :param batch_size: The size of batch
        :param seq_length: The length of sequence
        :return: Batches as a Numpy array
        """
        # TODO: Implement Function
        n_batches = int(len(int_text) / (batch_size * seq_length))

        # Drop the last few characters to make only full batches
        xdata = np.array(int_text[: n_batches * batch_size * seq_length])
        ydata = np.array(int_text[1: n_batches * batch_size * seq_length + 1])
        ydata[-1] = xdata[0]

        x_batches = np.split(xdata.reshape(batch_size, -1), n_batches, 1)
        y_batches = np.split(ydata.reshape(batch_size, -1), n_batches, 1)

        return np.array(list(zip(x_batches, y_batches)))

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_batches(get_batches)
```

Tests Passed

Neural Network Training

Hyperparameters

Tune the following parameters:

- Set num_epochs to the number of epochs.
- Set batch_size to the batch size.
- Set rnn_size to the size of the RNNs.
- Set embed_dim to the size of the embedding.
- Set seq_length to the length of sequence.
- Set learning_rate to the learning rate.
- Set show_every_n_batches to the number of batches the neural network should print progress.

```
In [14]: # Number of Epochs
num_epochs = 500
# Batch Size
batch_size = 100
# RNN Size
rnn_size = 500
# Embedding Dimension Size
embed_dim = 256
# Sequence Length
seq_length = 30
# Learning Rate
learning_rate = 0.001
# Show stats for every n number of batches
show_every_n_batches = 100

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

save_dir = './save'
```

Build the Graph

Build the graph using the neural network you implemented.

```
In [15]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

from tensorflow.contrib import seq2seq

train_graph = tf.Graph()
with train_graph.as_default():
    vocab_size = len(int_to_vocab)
    input_text, targets, lr = get_inputs()
    input_data_shape = tf.shape(input_text)
    cell, initial_state = get_init_cell(input_data_shape[0], rnn_size)
    logits, final_state = build_nn(cell, rnn_size, input_text, vocab_size,
    embed_dim)

    # Probabilities for generating words
    probs = tf.nn.softmax(logits, name='probs')

    # Loss function
    cost = seq2seq.sequence_loss(
        logits,
        targets,
        tf.ones([input_data_shape[0], input_data_shape[1]]))

    # Optimizer
    optimizer = tf.train.AdamOptimizer(lr)

    # Gradient Clipping
    gradients = optimizer.compute_gradients(cost)
    capped_gradients = [(tf.clip_by_value(grad, -1., 1.), var) for grad, va
r in gradients if grad is not None]
    train_op = optimizer.apply_gradients(capped_gradients)
```

Train

Train the neural network on the preprocessed data. If you have a hard time getting a good loss, check the [forums](https://discussions.udacity.com/) (<https://discussions.udacity.com/>) to see if anyone is having the same problem.

```

In [16]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

batches = get_batches(int_text, batch_size, seq_length)

with tf.Session(graph=train_graph) as sess:
    sess.run(tf.global_variables_initializer())

    for epoch_i in range(num_epochs):
        state = sess.run(initial_state, {input_text: batches[0][0]})

        for batch_i, (x, y) in enumerate(batches):
            feed = {
                input_text: x,
                targets: y,
                initial_state: state,
                lr: learning_rate}
            train_loss, state, _ = sess.run([cost, final_state, train_op],
            feed)

            # Show every <show_every_n_batches> batches
            if (epoch_i * len(batches) + batch_i) % show_every_n_batches ==
0:
                print('Epoch {:>3} Batch {:>4}/{}    train_loss = {:.3f}'.fo
rmat(
                    epoch_i,
                    batch_i,
                    len(batches),
                    train_loss))

        # Save Model
        saver = tf.train.Saver()
        saver.save(sess, save_dir)
        print('Model Trained and Saved')

```

Epoch	0	Batch	0/23	train_loss = 8.820
Epoch	4	Batch	8/23	train_loss = 5.406
Epoch	8	Batch	16/23	train_loss = 5.007
Epoch	13	Batch	1/23	train_loss = 4.508
Epoch	17	Batch	9/23	train_loss = 4.246
Epoch	21	Batch	17/23	train_loss = 4.037
Epoch	26	Batch	2/23	train_loss = 3.952
Epoch	30	Batch	10/23	train_loss = 3.786
Epoch	34	Batch	18/23	train_loss = 3.650
Epoch	39	Batch	3/23	train_loss = 3.482
Epoch	43	Batch	11/23	train_loss = 3.397
Epoch	47	Batch	19/23	train_loss = 3.240
Epoch	52	Batch	4/23	train_loss = 3.072
Epoch	56	Batch	12/23	train_loss = 2.926
Epoch	60	Batch	20/23	train_loss = 2.839
Epoch	65	Batch	5/23	train_loss = 2.807
Epoch	69	Batch	13/23	train_loss = 2.738
Epoch	73	Batch	21/23	train_loss = 2.630
Epoch	78	Batch	6/23	train_loss = 2.567
Epoch	82	Batch	14/23	train_loss = 2.506
Epoch	86	Batch	22/23	train_loss = 2.414
Epoch	91	Batch	7/23	train_loss = 2.364
Epoch	95	Batch	15/23	train_loss = 2.220
Epoch	100	Batch	0/23	train_loss = 2.199
Epoch	104	Batch	8/23	train_loss = 2.116
Epoch	108	Batch	16/23	train_loss = 2.028
Epoch	113	Batch	1/23	train_loss = 1.949
Epoch	117	Batch	9/23	train_loss = 1.906
Epoch	121	Batch	17/23	train_loss = 1.813
Epoch	126	Batch	2/23	train_loss = 1.774
Epoch	130	Batch	10/23	train_loss = 1.813
Epoch	134	Batch	18/23	train_loss = 1.692
Epoch	139	Batch	3/23	train_loss = 1.624
Epoch	143	Batch	11/23	train_loss = 1.602
Epoch	147	Batch	19/23	train_loss = 1.552
Epoch	152	Batch	4/23	train_loss = 1.513
Epoch	156	Batch	12/23	train_loss = 1.409
Epoch	160	Batch	20/23	train_loss = 1.399
Epoch	165	Batch	5/23	train_loss = 1.379
Epoch	169	Batch	13/23	train_loss = 1.321
Epoch	173	Batch	21/23	train_loss = 1.355
Epoch	178	Batch	6/23	train_loss = 1.297
Epoch	182	Batch	14/23	train_loss = 1.278
Epoch	186	Batch	22/23	train_loss = 1.206
Epoch	191	Batch	7/23	train_loss = 1.169
Epoch	195	Batch	15/23	train_loss = 1.068
Epoch	200	Batch	0/23	train_loss = 1.114
Epoch	204	Batch	8/23	train_loss = 1.097
Epoch	208	Batch	16/23	train_loss = 0.996
Epoch	213	Batch	1/23	train_loss = 0.990
Epoch	217	Batch	9/23	train_loss = 0.994
Epoch	221	Batch	17/23	train_loss = 0.912
Epoch	226	Batch	2/23	train_loss = 0.903
Epoch	230	Batch	10/23	train_loss = 0.926
Epoch	234	Batch	18/23	train_loss = 0.844
Epoch	239	Batch	3/23	train_loss = 0.797
Epoch	243	Batch	11/23	train_loss = 0.813

Epoch 247	Batch	19/23	train_loss = 0.775
Epoch 252	Batch	4/23	train_loss = 0.752
Epoch 256	Batch	12/23	train_loss = 0.709
Epoch 260	Batch	20/23	train_loss = 0.701
Epoch 265	Batch	5/23	train_loss = 0.664
Epoch 269	Batch	13/23	train_loss = 0.657
Epoch 273	Batch	21/23	train_loss = 0.686
Epoch 278	Batch	6/23	train_loss = 0.627
Epoch 282	Batch	14/23	train_loss = 0.642
Epoch 286	Batch	22/23	train_loss = 0.607
Epoch 291	Batch	7/23	train_loss = 0.596
Epoch 295	Batch	15/23	train_loss = 0.518
Epoch 300	Batch	0/23	train_loss = 0.551
Epoch 304	Batch	8/23	train_loss = 0.544
Epoch 308	Batch	16/23	train_loss = 0.482
Epoch 313	Batch	1/23	train_loss = 0.512
Epoch 317	Batch	9/23	train_loss = 0.499
Epoch 321	Batch	17/23	train_loss = 0.449
Epoch 326	Batch	2/23	train_loss = 0.464
Epoch 330	Batch	10/23	train_loss = 0.485
Epoch 334	Batch	18/23	train_loss = 0.435
Epoch 339	Batch	3/23	train_loss = 0.402
Epoch 343	Batch	11/23	train_loss = 0.408
Epoch 347	Batch	19/23	train_loss = 0.401
Epoch 352	Batch	4/23	train_loss = 0.371
Epoch 356	Batch	12/23	train_loss = 0.372
Epoch 360	Batch	20/23	train_loss = 0.356
Epoch 365	Batch	5/23	train_loss = 0.353
Epoch 369	Batch	13/23	train_loss = 0.332
Epoch 373	Batch	21/23	train_loss = 0.351
Epoch 378	Batch	6/23	train_loss = 0.325
Epoch 382	Batch	14/23	train_loss = 0.336
Epoch 386	Batch	22/23	train_loss = 0.331
Epoch 391	Batch	7/23	train_loss = 0.322
Epoch 395	Batch	15/23	train_loss = 0.294
Epoch 400	Batch	0/23	train_loss = 0.295
Epoch 404	Batch	8/23	train_loss = 0.295
Epoch 408	Batch	16/23	train_loss = 0.267
Epoch 413	Batch	1/23	train_loss = 0.265
Epoch 417	Batch	9/23	train_loss = 0.274
Epoch 421	Batch	17/23	train_loss = 0.241
Epoch 426	Batch	2/23	train_loss = 0.250
Epoch 430	Batch	10/23	train_loss = 0.278
Epoch 434	Batch	18/23	train_loss = 0.245
Epoch 439	Batch	3/23	train_loss = 0.245
Epoch 443	Batch	11/23	train_loss = 0.248
Epoch 447	Batch	19/23	train_loss = 0.235
Epoch 452	Batch	4/23	train_loss = 0.244
Epoch 456	Batch	12/23	train_loss = 0.218
Epoch 460	Batch	20/23	train_loss = 0.208
Epoch 465	Batch	5/23	train_loss = 0.219
Epoch 469	Batch	13/23	train_loss = 0.217
Epoch 473	Batch	21/23	train_loss = 0.217
Epoch 478	Batch	6/23	train_loss = 0.203
Epoch 482	Batch	14/23	train_loss = 0.222
Epoch 486	Batch	22/23	train_loss = 0.212
Epoch 491	Batch	7/23	train_loss = 0.210

Epoch 495 Batch 15/23 train_loss = 0.185
Model Trained and Saved

Save Parameters

Save seq_length and save_dir for generating a new TV script.

```
In [17]: """  
DON'T MODIFY ANYTHING IN THIS CELL  
"""  
  
# Save parameters for checkpoint  
helper.save_params((seq_length, save_dir))
```

Checkpoint

```
In [18]: """  
DON'T MODIFY ANYTHING IN THIS CELL  
"""  
  
import tensorflow as tf  
import numpy as np  
import helper  
import problem_unittests as tests  
  
_, vocab_to_int, int_to_vocab, token_dict = helper.load_preprocess()  
seq_length, load_dir = helper.load_params()
```

Implement Generate Functions

Get Tensors

Get tensors from loaded_graph using the function `get_tensor_by_name()`. (https://www.tensorflow.org/api_docs/python/tf/Graph#get_tensor_by_name). Get the tensors using the following names:

- "input:0"
- "initial_state:0"
- "final_state:0"
- "probs:0"

Return the tensors in the following tuple (InputTensor, InitialStateTensor, FinalStateTensor, ProbsTensor)

```
In [19]: def get_tensors(loaded_graph):
        """
        Get input, initial state, final state, and probabilities tensor from <loaded_graph>
        :param loaded_graph: TensorFlow graph loaded from file
        :return: Tuple (InputTensor, InitialStateTensor, FinalStateTensor, ProbsTensor)
        """
        # TODO: Implement Function
        InputTensor = loaded_graph.get_tensor_by_name("input:0")
        InitialStateTensor = loaded_graph.get_tensor_by_name("initial_state:0")
        FinalStateTensor = loaded_graph.get_tensor_by_name("final_state:0")
        ProbsTensor = loaded_graph.get_tensor_by_name("probs:0")
        return InputTensor, InitialStateTensor, FinalStateTensor, ProbsTensor

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_tensors(get_tensors)
```

Tests Passed

Choose Word

Implement the `pick_word()` function to select the next word using probabilities.

```
In [20]: def pick_word(probabilities, int_to_vocab):
        """
        Pick the next word in the generated text
        :param probabilities: Probabilities of the next word
        :param int_to_vocab: Dictionary of word ids as the keys and words as the values
        :return: String of the predicted word
        """
        # TODO: Implement Function
        word = np.random.choice(list(int_to_vocab.values()), p=probabilities)
        return word

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_pick_word(pick_word)
```

Tests Passed

Generate TV Script

This will generate the TV script for you. Set `gen_length` to the length of TV script you want to generate.

```

In [21]: gen_length = 200
         # homer_simpson, moe_szyslak, or Barney_Gumble
         prime_word = 'moe_szyslak'

         """
         DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
         """

         loaded_graph = tf.Graph()
         with tf.Session(graph=loaded_graph) as sess:
             # Load saved model
             loader = tf.train.import_meta_graph(load_dir + '.meta')
             loader.restore(sess, load_dir)

             # Get Tensors from Loaded model
             input_text, initial_state, final_state, probs = get_tensors(loaded_graph)

             # Sentences generation setup
             gen_sentences = [prime_word + ':']
             prev_state = sess.run(initial_state, {input_text: np.array([[1]])})

             # Generate sentences
             for n in range(gen_length):
                 # Dynamic Input
                 dyn_input = [[vocab_to_int[word] for word in gen_sentences[-seq_length:]]]

                 dyn_seq_length = len(dyn_input[0])

                 # Get Prediction
                 probabilities, prev_state = sess.run(
                     [probs, final_state],
                     {input_text: dyn_input, initial_state: prev_state})

                 pred_word = pick_word(probabilities[dyn_seq_length-1], int_to_vocab)

                 gen_sentences.append(pred_word)

             # Remove tokens
             tv_script = ' '.join(gen_sentences)
             for key, token in token_dict.items():
                 ending = ' ' if key in ['\n', '(', "'"] else ''
                 tv_script = tv_script.replace(' ' + token.lower(), key)
             tv_script = tv_script.replace('\n ', '\n')
             tv_script = tv_script.replace('(', '(')

             print(tv_script)

```

```
moe_szyslak: gee, i can't--
barney_gumble: hey!, i've been even come from" face and" and" bonfire"?
moe_szyslak: hey, barney. my name put on the way.
seymour_skinner: we got caught in this special?
flea:(aside) and you all hate that last night / bartender he doesn't stand th
ere and be able to bring her out!
homer_simpson:(gasp)" african princess."
gil_gunderson:(low voice) really?
marge_simpson: gimme that on the nose...
moe_szyslak:(sings) good king wenceslas looked out on their mouth."
moe_szyslak:(embarrassed) whoa, tha...(laughs) that, no. you're a fabulous ca
tch. see the day of speech never need your name, and on the bar, then you're
just tryin' to.
homer_simpson: well, you really think i could do these with marge?
lenny_leonard: quiet, we're great.
moe_szyslak: look, i really a little trick. i pictured everyone in their unde
rwear. the judge,
```

The TV Script is Nonsensical

It's ok if the TV script doesn't make any sense. We trained on less than a megabyte of text. In order to get good results, you'll have to use a smaller vocabulary or get more data. Luckily there's more data! As we mentioned in the begging of this project, this is a subset of [another dataset \(https://www.kaggle.com/wcukierski/the-simpsons-by-the-data\)](https://www.kaggle.com/wcukierski/the-simpsons-by-the-data). We didn't have you train on all the data, because that would take too long. However, you are free to train your neural network on all the data. After you complete the project, of course.

Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd_tv_script_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "helper.py" and "problem_unittests.py" files in your submission.